

# CS 149 Operating Systems

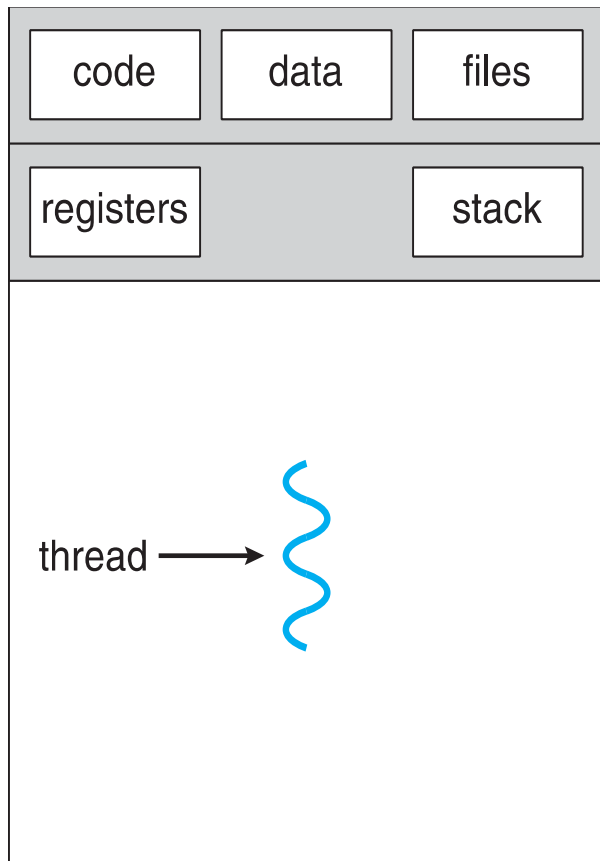
## ***Threads***

Instructor: Kong Li

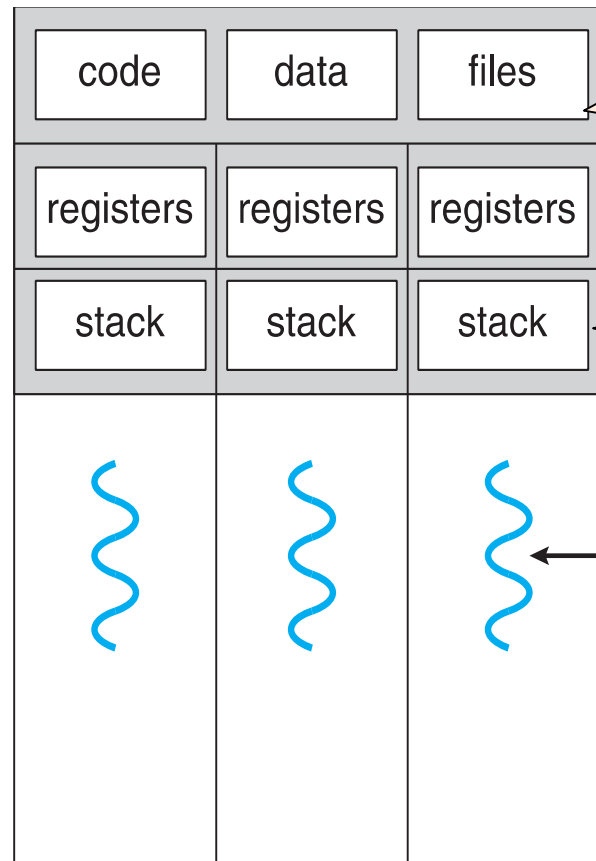
# Content

- What & benefits
- Multicore Programming
- Thread Programming
- Thread Libraries
  - Pthread
  - Java
- Semantics: fork, exec, signal

# Single and Multithreaded Processes



single-threaded process




multithreaded process

Threads share entire memory space of proc (code, data, files, etc)

Each thread has own set of registers, PC, SP, etc.


Tradeoffs?  
Pros? Cons?

# Thread Pros and Cons


- Advantages of using threads:
    - Responsiveness
      - allow continued execution if part of proc is blocked
      - important for user interfaces
    - Resource sharing
      - easier than shared memory or message passing
    - Economy
      - cheaper than proc creation
      - overhead(thread switching) < overhead(proc context switching)
    - Scalability
      - take advantage of multiprocessor architectures
  - Disadvantages of using threads
    - Synchronization overhead
    - Many library functions not thread-safe. E.g., random vs random\_r
    - Lack of robustness (one thread can impact all threads)
- 

# Concurrency vs. Parallelism

- **Parallelism:** perform more than one task **simultaneously**
- **Concurrency:** more than one task making progress
  - Either **interleaved or in parallel or hybrid**
  - Concurrency w/o parallelism is possible
  - Single processor / core: scheduler providing concurrency
- Concurrent execution on single-core system:



- Parallelism on a multi-core system:

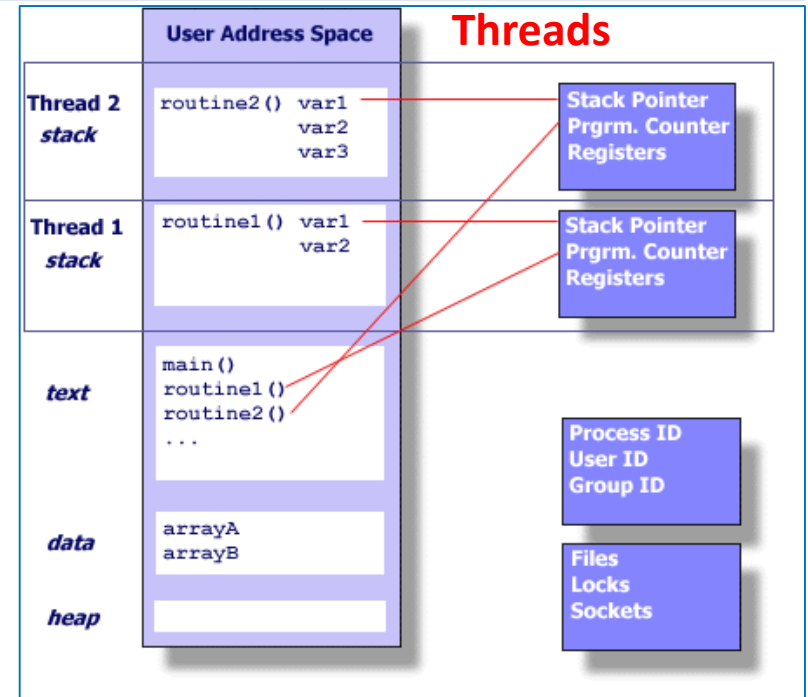
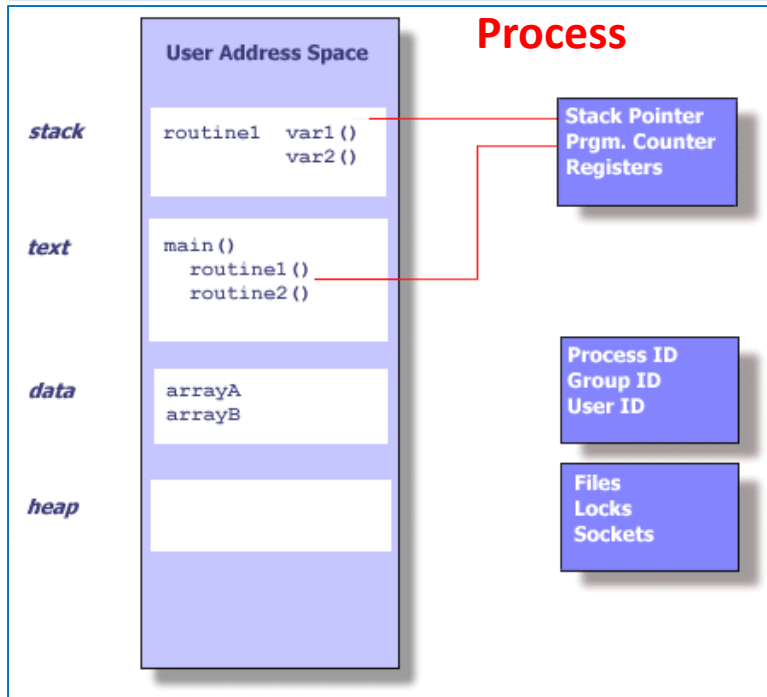


- S: serial portion of app
- As  $N \rightarrow \infty$ , speedup  $\rightarrow 1/S$
- App's serial portion becomes the dominating factor for performance

- Amdahl's Law

$$\text{speedup} \leq \frac{1}{S + \frac{(1-S)}{N}}$$

# Proc vs Thread

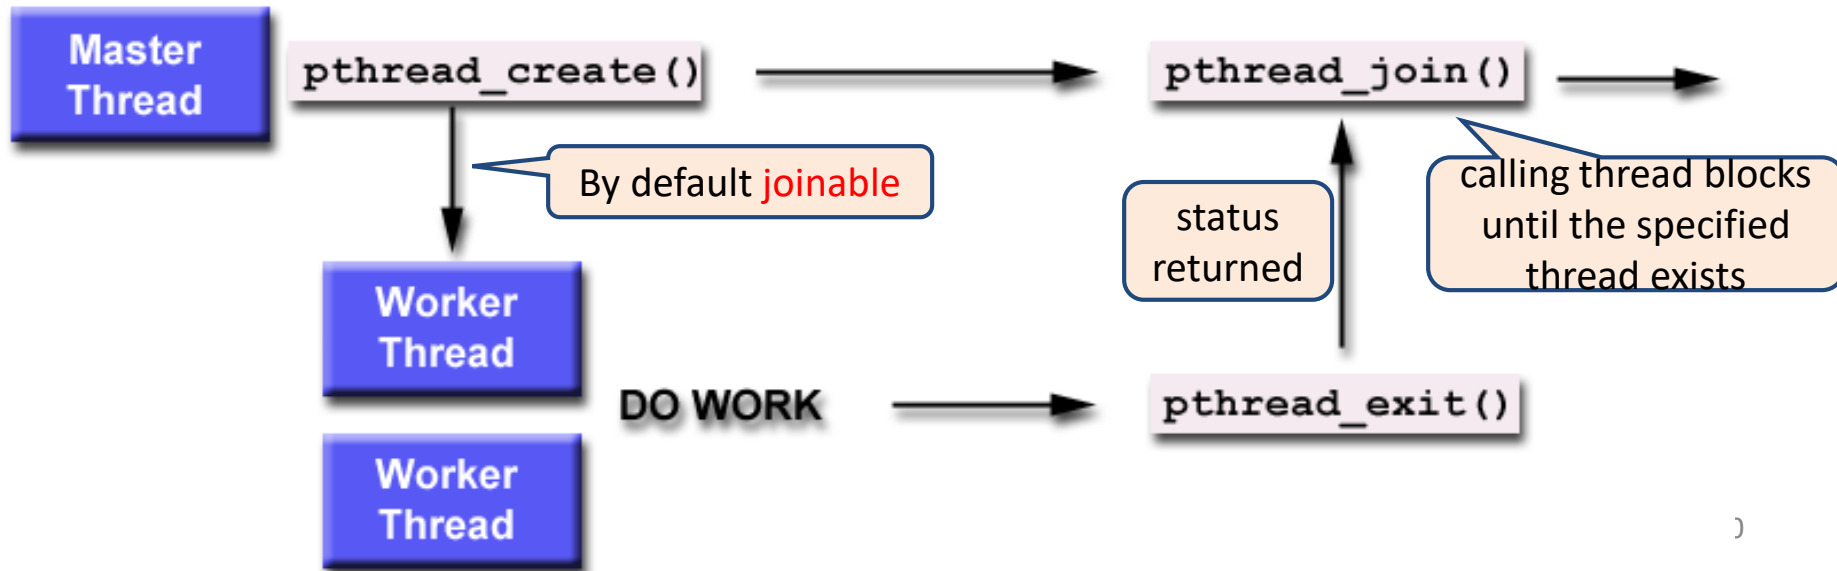
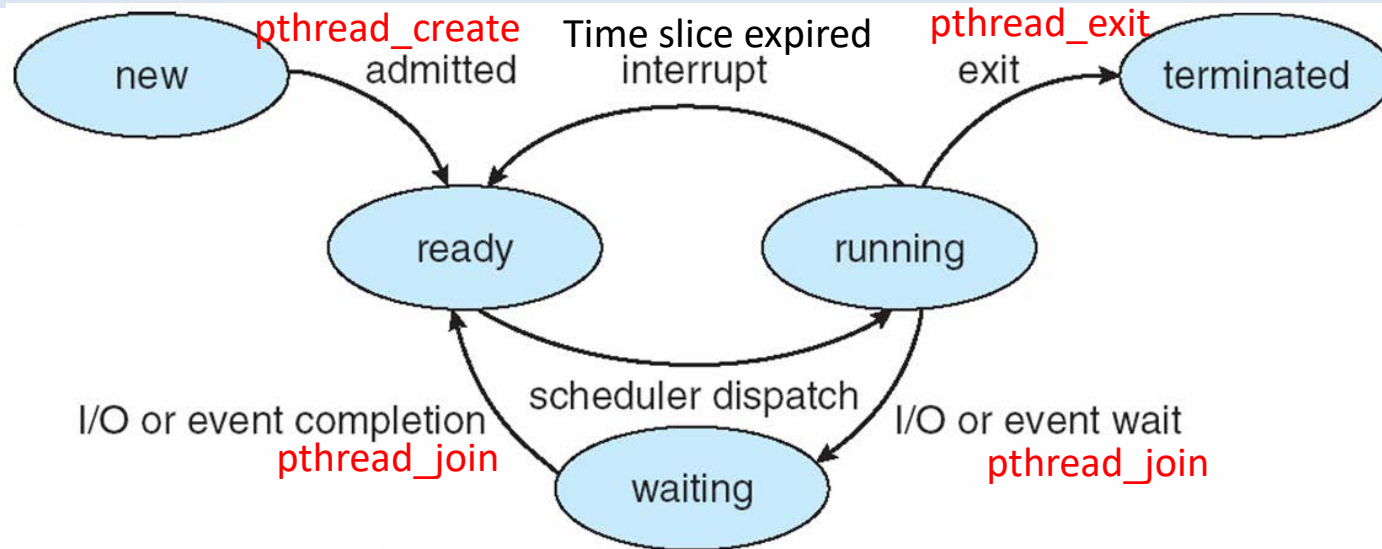


- Threads in one proc
  - Share: heap, data, code (text), global variables, file descriptors, proc id, etc.
  - Individual (TCB): stack, registers, program counter, thread id, etc.
- Some system calls affect only thread (e.g., `_r`), some affect the entire proc (e.g., `exit`)
- Linux API `_r`: **thread safe**. E.g., `random` vs `random_r`
- Linux: **errno** – thread-local (one instance per thread)

# Thread library: Pthreads

- POSIX standard (IEEE 1003.1c) API
  - Specification, not implementation
  - API specifies behavior of thread library
  - implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)
- `man pthreads`
- Compilation on Linux: `gcc test.c -pthread`
  - Not recommended: `gcc test.c -lpthread`
- Eclipse IDE: by default does not include “-pthread” in project settings
  - Project Explorer pane: project -> Properties -> C/C++ build -> Settings -> Tool Settings ->
    - GCC C Compiler -> Miscellaneous, Add “-pthread” into the beginning of “Other Flags”
    - GCC C Linker -> Miscellaneous, Add “-pthread” into the beginning of “Linker Flags”
- <https://computing.llnl.gov/tutorials/pthreads/>

# Thread State and Lifecycle





# Thread Creation

```
#include <pthread.h>
int pthread_create(pthread_t *thread,
                  const pthread_attr_t *attr,
                  void*(*start_routine) (void *), void *arg);
```

- Create a new thread which starts executing `start_routine`, with the parameter `arg`
- init its attrs using `attr`
- thread attributes:
  - priority, stack size, name, detach-state, etc.
- by default: **Joinable thread**
- Q: why **detached** thread?

## Linux thread attributes:

Detach state = **PTHREAD\_CREATE\_JOINABLE**

Scope = **PTHREAD\_SCOPE\_SYSTEM**

Inherit scheduler = **PTHREAD\_INHERIT\_SCHED**

Scheduling policy = **SCHED\_OTHER**

Scheduling priority = 0

Guard size = 4096 bytes

Stack address = 0x40196000

Stack size = 0x201000 bytes

# Thread Termination & Join

```
#include <pthread.h>
```

```
void pthread_exit(void *value_ptr);
```

```
int pthread_join(pthread_t thread, void **value_ptr);
```

should **not** be located on the calling thread's stack (**why?**)

- `pthread_exit()`
  - Calling thread exits & returns status
- `pthread_join()`
  - Calling thread blocks until the specified thread exits
  - Get status back
- Thread termination
  - returns from its starting routine
  - call `pthread_exit()`
  - cancelled by `pthread_cancel()`
  - process terminated

Process	Thread
<code>fork()</code>	<code>pthread_create()</code>
<code>exit()</code>	<code>pthread_exit()</code>
<code>exec()</code>	N/A
<code>waitpid(pid, ...)</code>	<code>pthread_join()</code>
<code>wait()</code>	N/A
zombie process	zombie thread
<code>kill()</code>	<code>pthread_cancel()</code>
<code>getpid()</code>	<code>pthread_self()</code>

# Pthreads Example

gcc ... -pthread

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

```
/* get the default attributes */
pthread_attr_t attr;
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid, &attr, runner, argv[1]);
/* wait for the thread to exit */
pthread_join(tid, NULL);

printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

**Demo!**

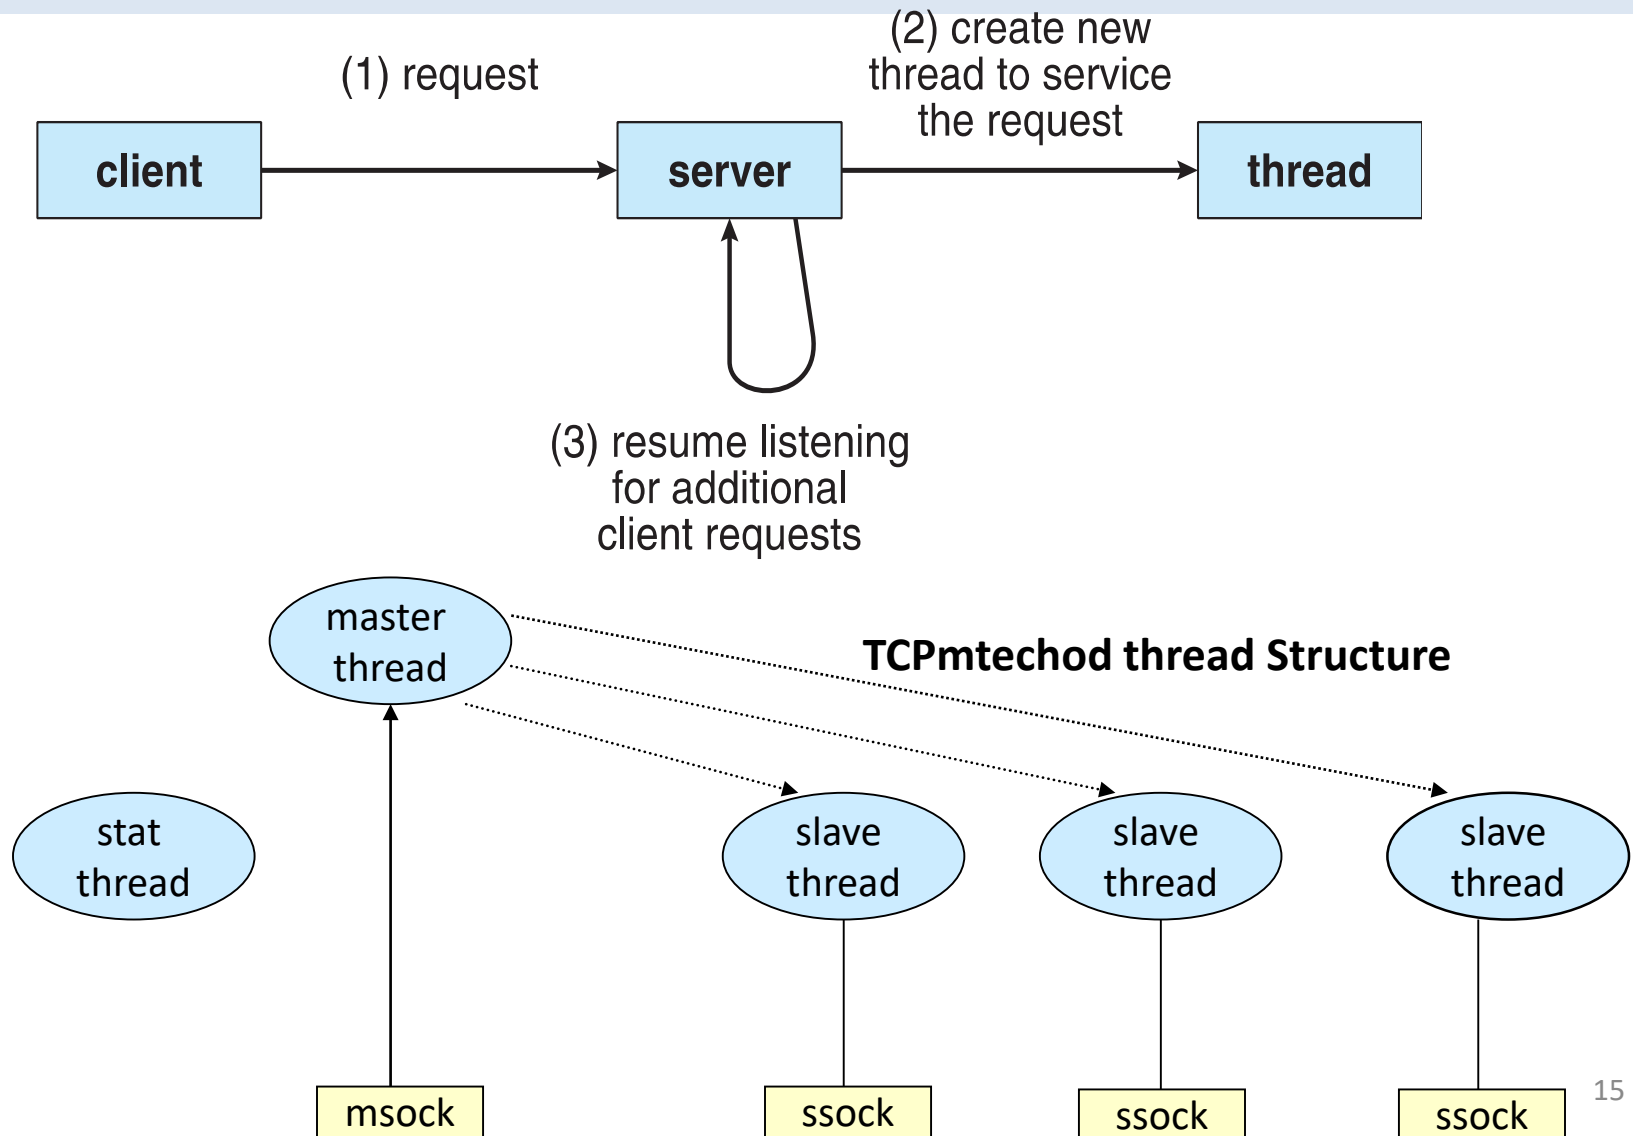
# Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

# Multithreaded Server Architecture



# Pthreads: TCPmtechod.c (Master Thread)

```
pthread_attr_t ta;
```

thread attributes obj

```
msock = passiveTCP(service, QLEN);
```

```
(void) pthread_attr_init(&ta);
```

```
(void) pthread_attr_setdetachstate(&ta, PTHREAD_CREATE_DETACHED);
```

```
(void) pthread_mutex_init(&stats.st_mutex, 0);
```

```
if (pthread_create(&tid, &ta, (void * (*)(void *))prstats, 0) < 0)
```

**detached thread**

- not joinable
- exit status discarded
- default is joinable

one stat thread

```
    errexit("pthread_create(prstats): %s\n", strerror(errno));
```

```
while (1) {
```

```
    alen = sizeof(fsin);
```

```
    ssock = accept(msock, (struct sockaddr *)&fsin, &alen);
```

```
    if (ssock < 0) {
```

```
        if (errno == EINTR)
```

```
            continue;
```

```
        errexit("accept: %s\n", strerror(errno));
```

```
    }
```

```
    if (pthread_create(&tid, &ta, (void * (*)(void *))TCPechnod, (void *)ssock) < 0)
```

```
        errexit("pthread_create: %s\n", strerror(errno));
```

```
}
```

Error code in errno

starting  
routine

parameter

many echo  
threads

# Thread Cancellation

- **Request to cancel** thread: terminating a thread before it has finished

```
int pthread_cancel(thread_t tid);
```

- Actual cancellation depends on thread state `pthread_setcancelstate`

default →

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

Why?

- If state == disabled, cancellation remains pending until thread enables it
- If state == enabled, determined by type `pthread_setcanceltype`
  - If type == async, terminates the target thread immediately
  - If type == deferred, cancellation occurs when thread reaches **cancellation point** (certain functions)
- Request delivery of pending cancel req: `pthread_testcancel()` ;
- Cleanup: invoke **cleanup handler(s)** async before terminating thread
- On Linux systems, thread cancellation is handled through **signals**

# Java Threads

- managed by the JVM
  - using the threads model provided by underlying OS
  - Creation:
    - Define a new class that
      - Derived from the Thread class
      - Override run() method
    - Define a new class that implements the Runnable interface
- ```
public interface Runnable
{
    public abstract void run();
}
```
- run( ): thread starting routine
  - Thread mgmt: java.util.concurrent package



# Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}
```

```
class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

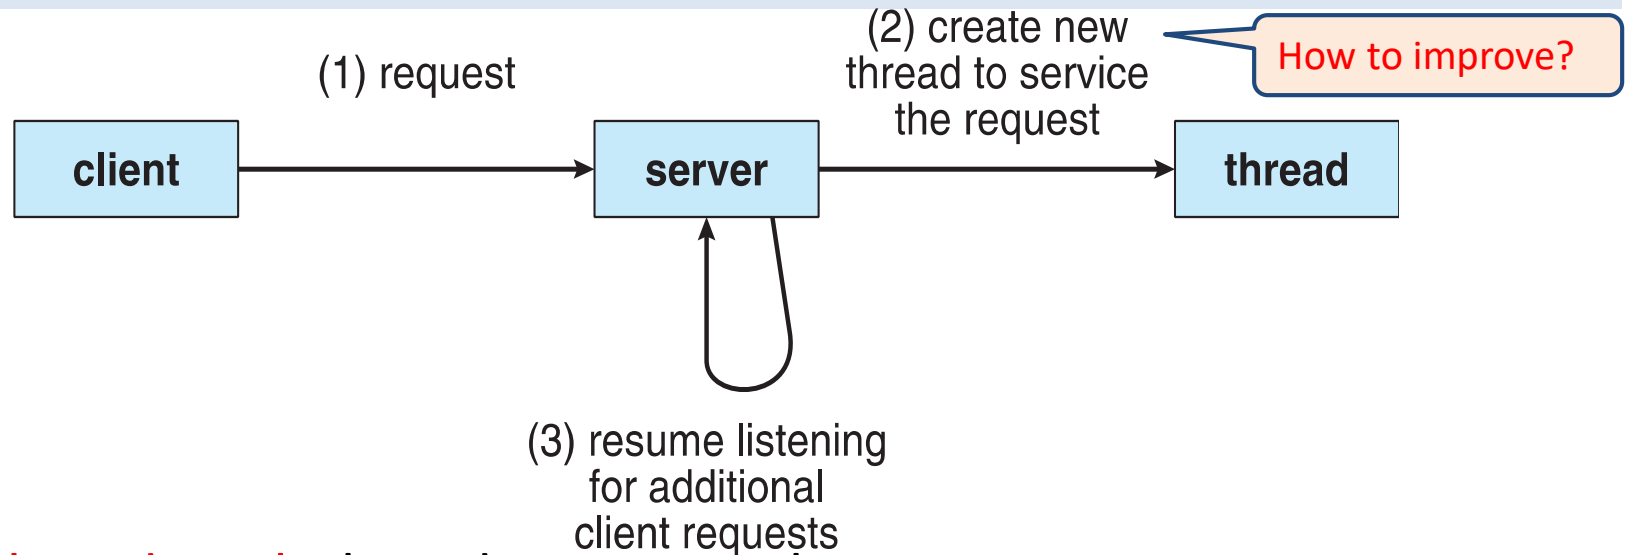
    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```

thread starting  
routine

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```

Demo!

# Multithreaded Server Architecture



- **thread pool**: threads await work
- Advantages:
  - (for server) **faster to service request than creating a new thread**
  - Allows # of threads in app(s) to be bound to the size of the pool
  - Separating task (to be performed) from mechanics (of creating task): different strategies for running task
    - Tasks could be scheduled to run periodically

# Semantics of `fork()` and `exec()`

- `fork()`: duplicate only the calling thread or all threads?
  - Some UNIXes have two versions of `fork`
  - Linux: dup **only** the calling thread in the new proc;  
**no other threads**
    - Cleanup other threads? `pthread_atfork()`
- `exec()`: **replace the running proc (including all threads)**

# Semantics of Signal Handling

- **Signals**: UNIX way to notify a proc that an event has occurred
  - Signal generated by particular event (inside or outside of a proc)
  - Signal delivered to a proc
  - Signal handling (**per-signal basis**): default action, ignore, catch
- man 7 signal
- Single-threaded: signal delivered to **process**
  - man 2 kill
  - API: `kill(pid_t pid, int signal);`
- Multi-threaded: signal delivered to which thread?
  - Thread specific: to thread causing the signal, e.g., SIGSEGV, SIGFPE
  - Proc specific: to any thread not blocking the signal, e.g., Control-C
  - API: `pthread_cancel(pthread_t tid, int signal);`

# Thread-Local Storage

- **Thread-local storage (TLS)**
  - Each thread to have its **own copy** of data; **unique to each thread**
  - Similar to static data but per thread basis. e.g., **errno**
  - When is it useful?
    - when one does not control thread creation (i.e., thread pool)
- **local variables vs TLS**
  - Local variables: visible only during single function invocation
  - TLS: visible across function invocations

# Summary

- Thread vs Process
  - Share what?
  - Which info is private to each thread (not shared with process)?
  - What is the content in thread control block (TCB)?
- Thread life cycle
  - Which state transition is triggered by which pthread API?
  - Pthread: create, termination, join, cancel
  - Process APIs vs Pthread APIs
  - Java thread
  - Multi-threaded server & Thread pool
- Fork vs exec
- Thread signal handling
- Thread local storage

# Self Exercises

- 4.7 ~ 4.9, 4.13, 4.15, 4.17
- 4.20, 4.22, 4.25